




Apache Cassandra[™] Fasttrack




Data Modeling



Data Modeling Overview





Conceptual Data Model



Purpose of Conceptual Model

- Understand your data
- Essential objects
- Constraints

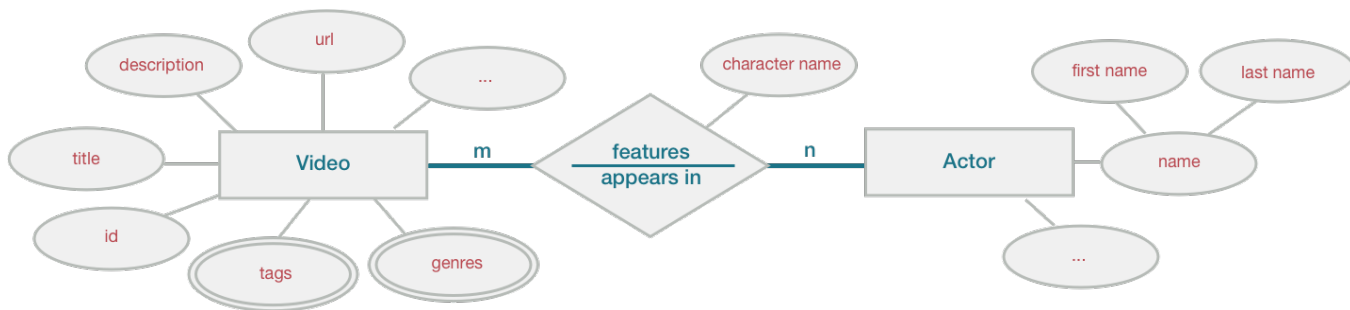


Collaboration = Key



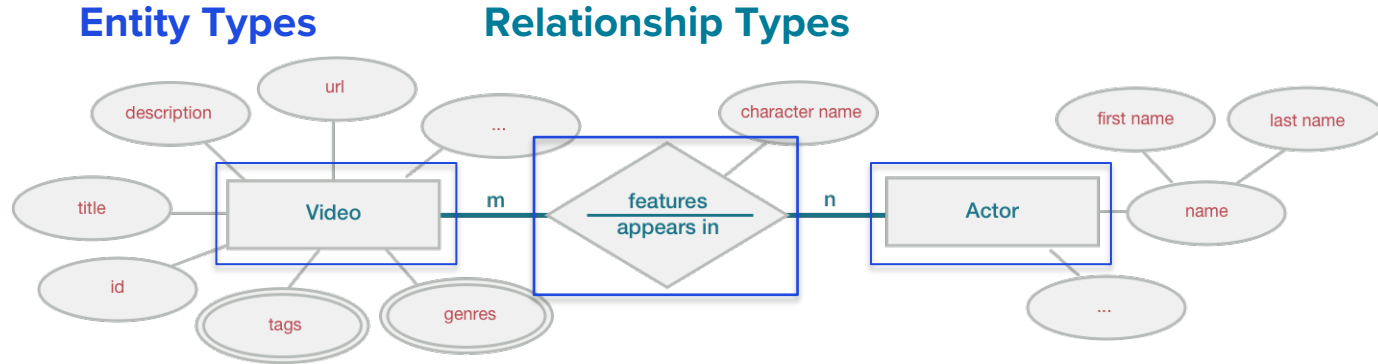
So What Does a Conceptual Model Look Like?

- Abstract view of your domain
- Technology independent
- Not specific to any database system



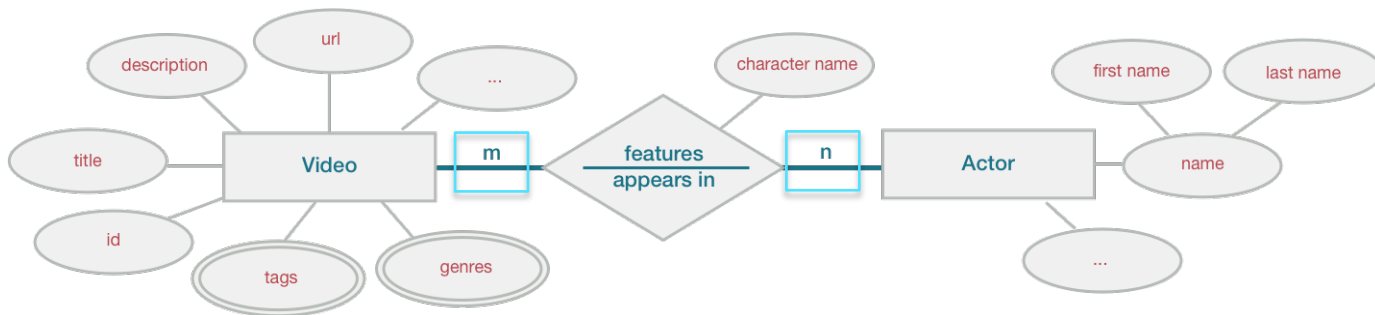
Entity-Relationship (ER) Model

- Entity Types - Relationship Types - Attribute Types



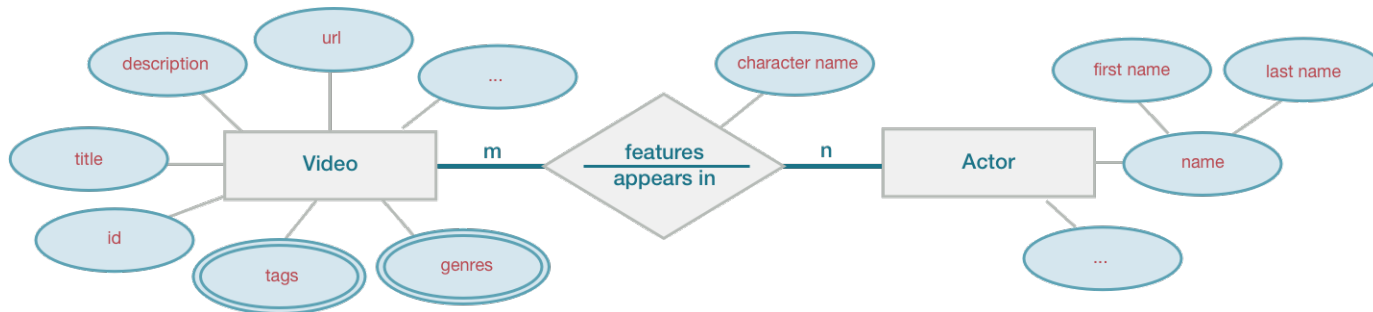
Cardinality

- Number of times an entity can/must participate in the relationship
- Other possibilities:
 - 1-n
 - 1-1



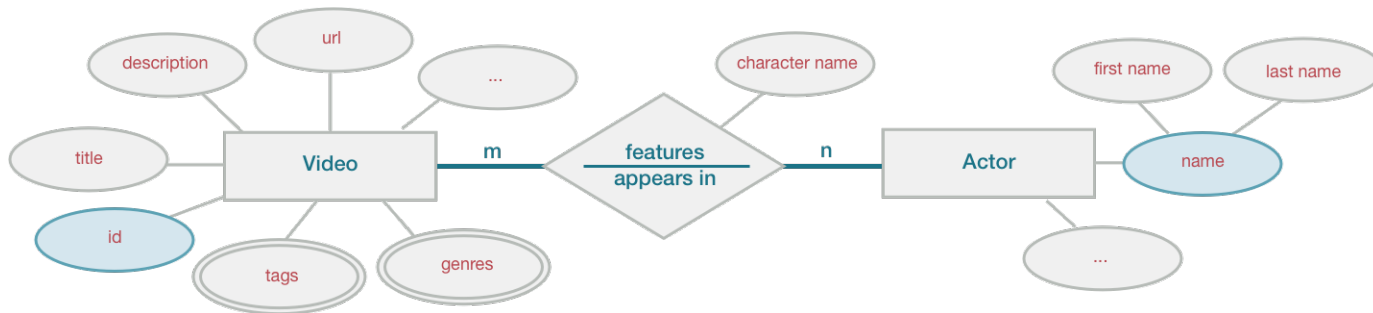
Attribute Types

- Fields to store data about an entity or relationship



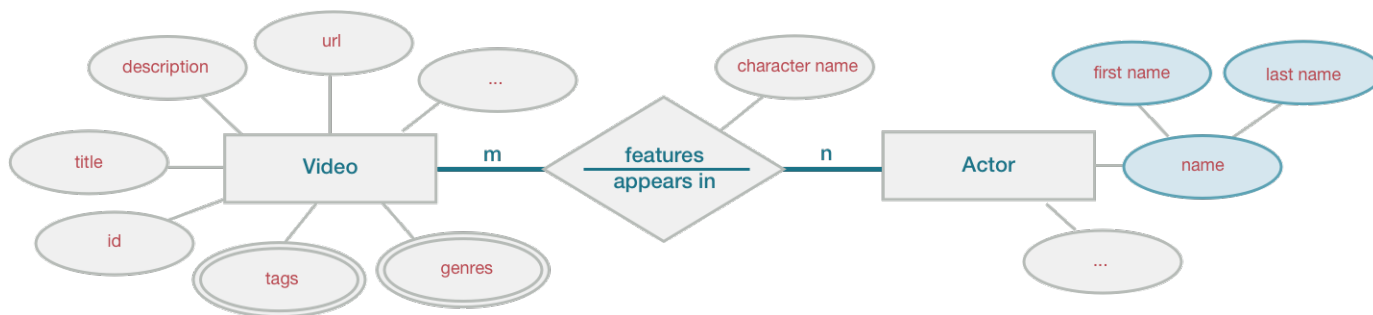
Key Attributes

- Identifies an Object



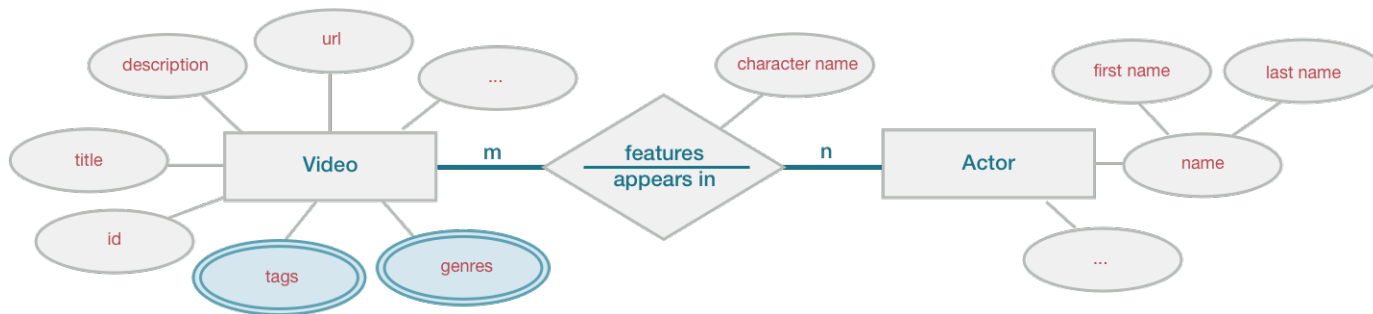
Composite Attributes

- Groups related attributes together

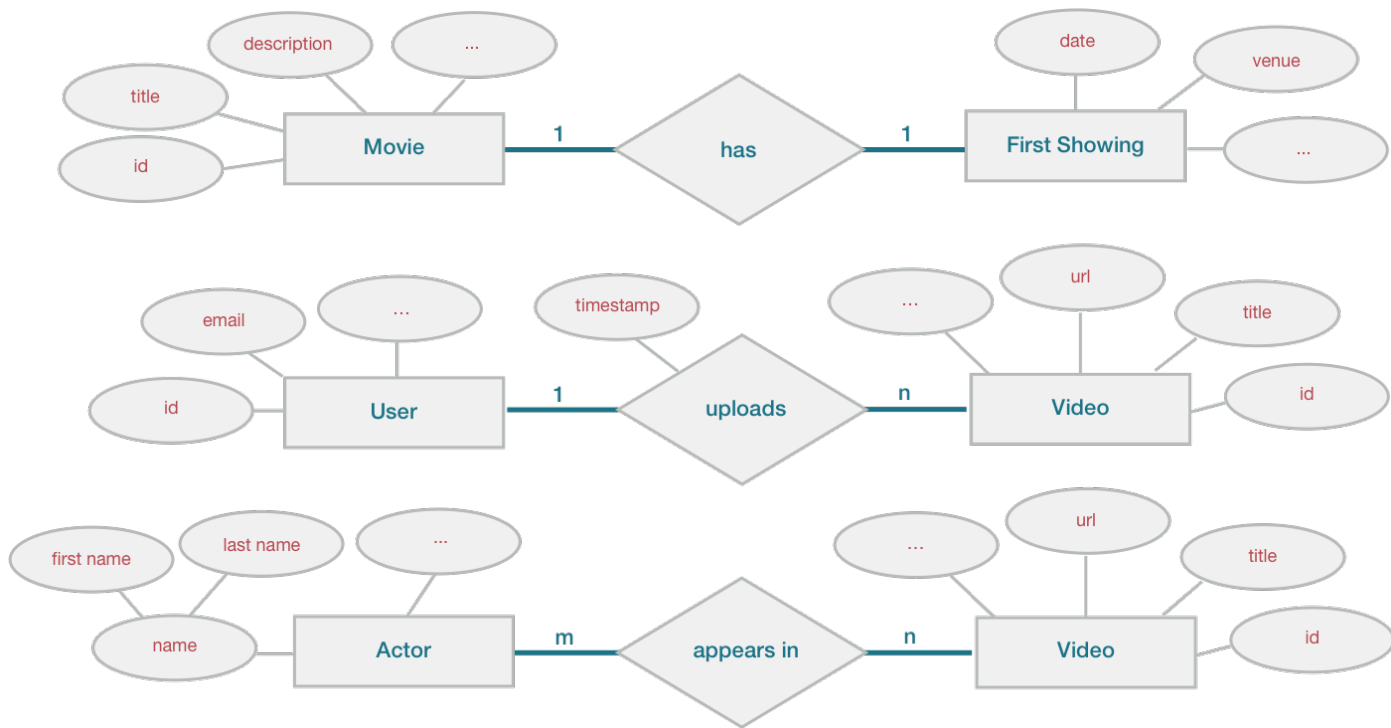


Multi Valued Attributes

- Stored multiple values per attribute




Relationship Keys



Weak Entity Types


- Entities that cannot exist without a String entity type



Exercise 03.01




Finish the Conceptual Model




Exercise 03.01 – Finish the Conceptual Model

- Model the KillrVideo entities
- Identify the entity types, relationship types, and attribute types

Workflow and Access Patterns




Where Are We Now?



Workflow

- Each application has a workflow—Tasks/causal dependencies form a graph
- Access patterns help determine how data is accessed—Know what queries you will run first
- **Example Task:**
 - Have a user login to a site

Workflow & Access Patterns



ACCESS PATTERNS

Q1: Find a user with a **specified email**

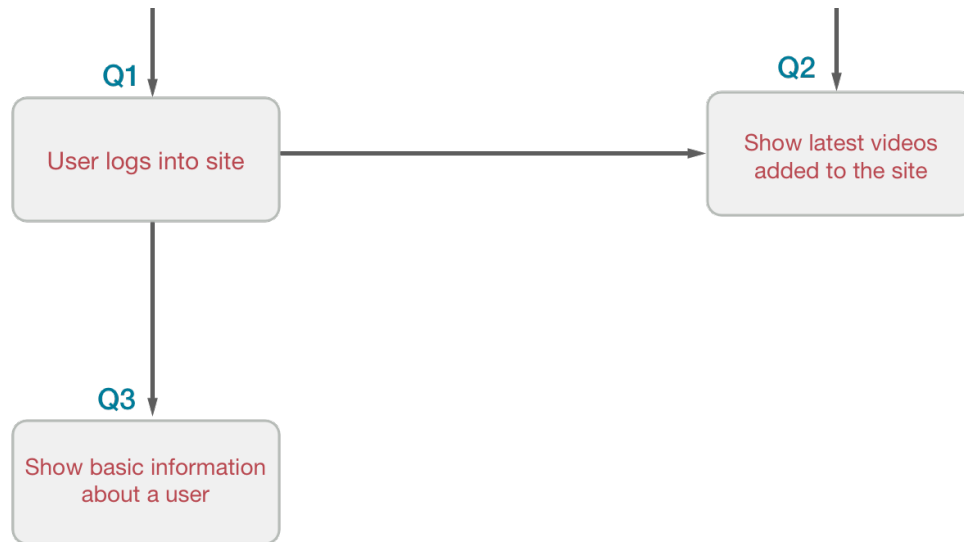
Workflow & Access Patterns



ACCESS PATTERNS

Q1: Find a user with a **specified email**
Q2: Find most recently uploaded **videos**

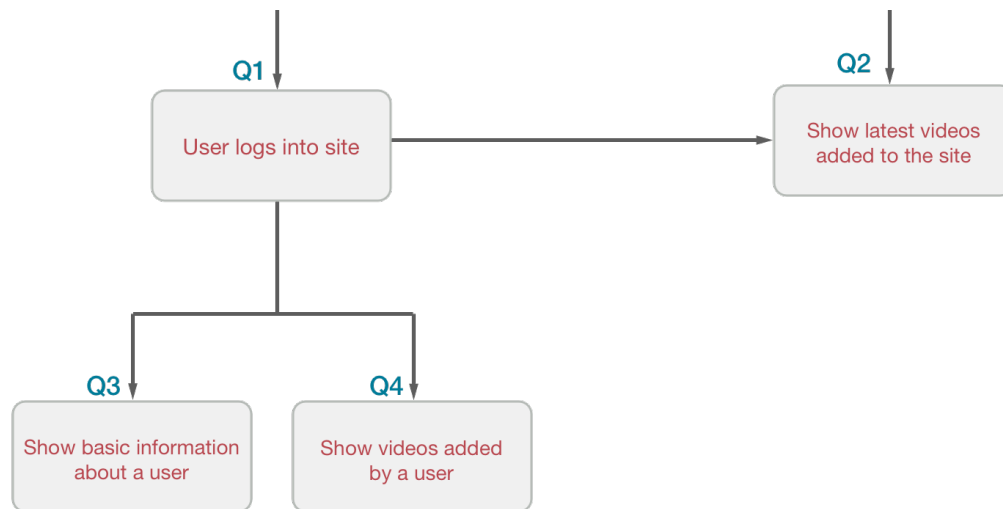
Workflow & Access Patterns



ACCESS PATTERNS

- Q1: Find a user with a **specified email**
- Q2: Find most recently uploaded **videos**
- Q3: Find a **user with a specified id**

Workflow & Access Patterns



ACCESS PATTERNS

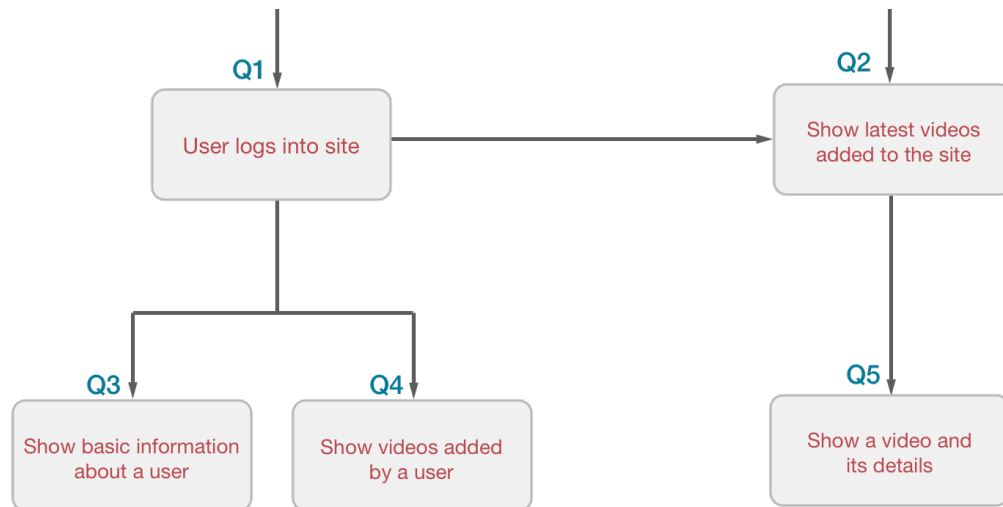
Q1: Find a user with a **specified email**

Q2: Find most recently uploaded **videos**

Q3: Find a **user with a specified id**

Q4: Find videos uploaded by a **user with a known id**(show most recently uploaded videos first)

Workflow & Access Patterns



ACCESS PATTERNS

Q1: Find a user with a **specified email**

Q2: Find most recently uploaded **videos**

Q3: Find a **user with a specified id**


Q4: Find videos uploaded by a **user with a known id**(show most recently uploaded videos first)

Q5: Find a video with a **specified video id**

Exercise 03.02




Finish the Workflow & Access Patterns




Exercise 03.02 – Finish the Workflow & Access Patterns

- Finish the application workflow based KillrVideo query requirements
- Add the remaining access patterns based on KillrVideo query requirements

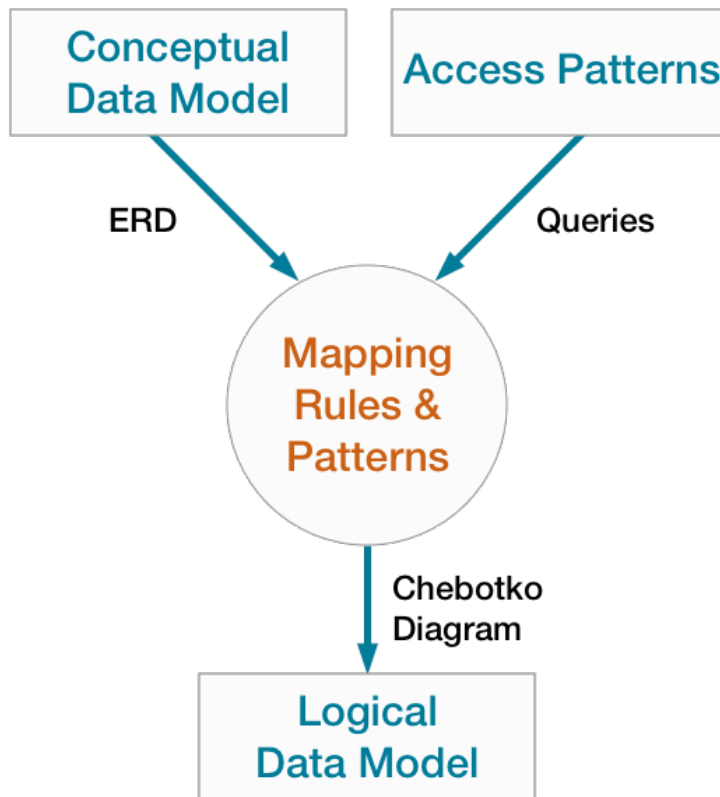
Mapping Conceptual to Logical



Where are we now?

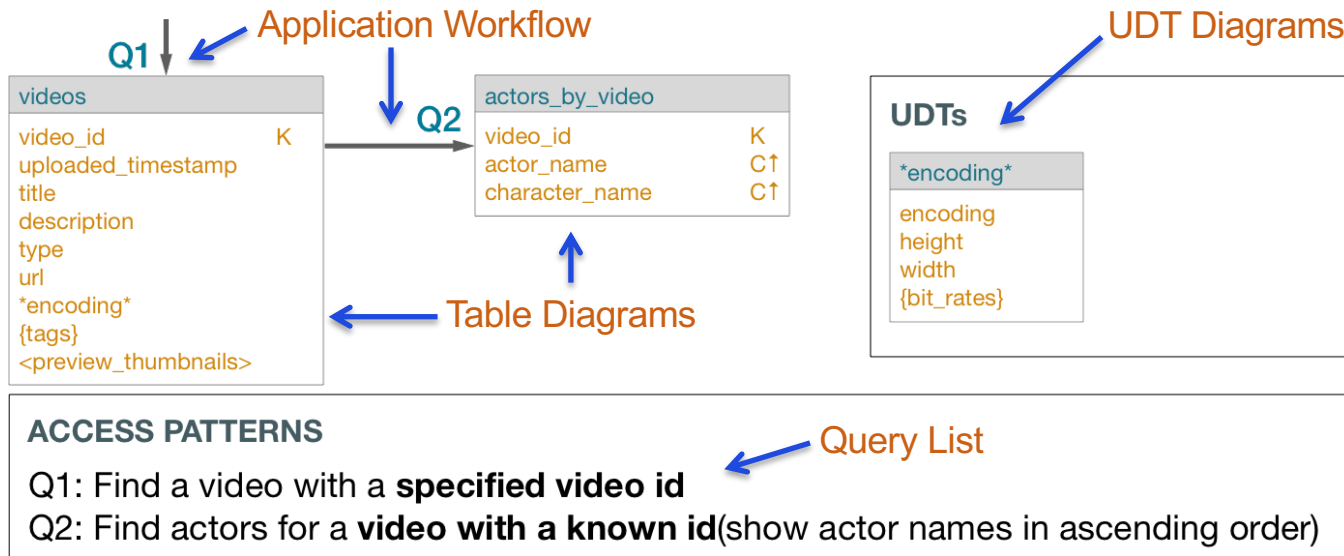


Query Driven Data Modeling



Chebotko Diagrams

- Graphical representation of Apache Cassandra™ database schema design
- Documents the logical and physical data model



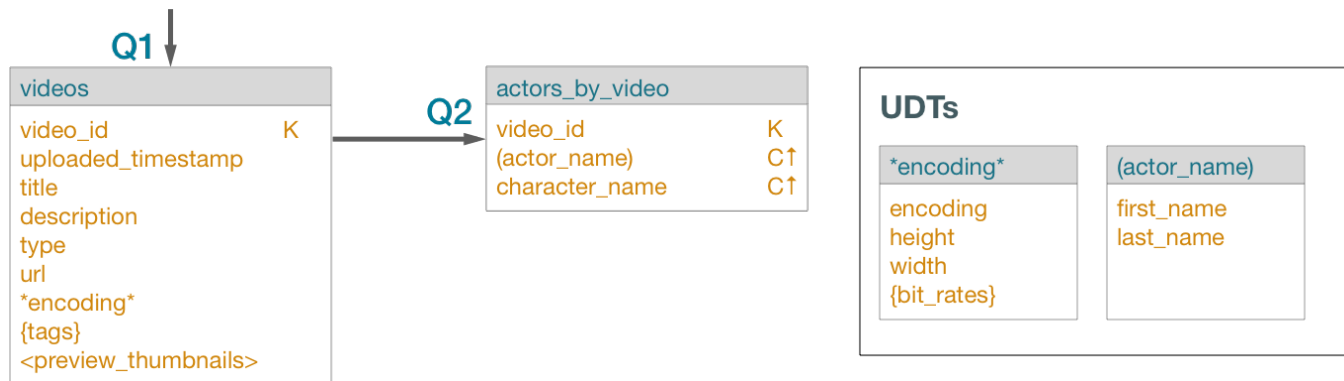
Chebotko Diagram Notations

- Logical-level shows column names and properties
- Physical-level also shows the column data type

table_name		
column_name_1	CQL Type K	← Partition key column
column_name_2	CQL Type C↑	← Clustering key column (ASC)
column_name_3	CQL Type C↓	← Clustering key column (DESC)
column_name_4	CQL Type S	← Static column
column_name_5	CQL Type IDX	← Secondary index column
column_name_6	CQL Type ++	← Counter column
[column_name_7]	CQL Type	← Collection column (list)
{column_name_8}	CQL Type	← Collection column (set)
<column_name_9>	CQL Type	← Collection column (map)
column_name_10	UDT Name	← UDT column
(column_name_11)	CQL Type	← Tuple column
column_name_12	CQL Type	← Regular column

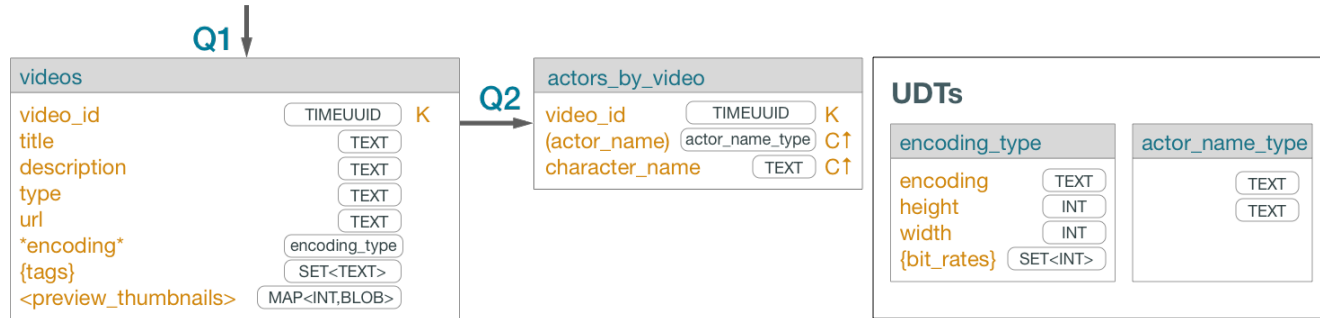
Logical UDT Diagram

- Represents user defined types and tuples

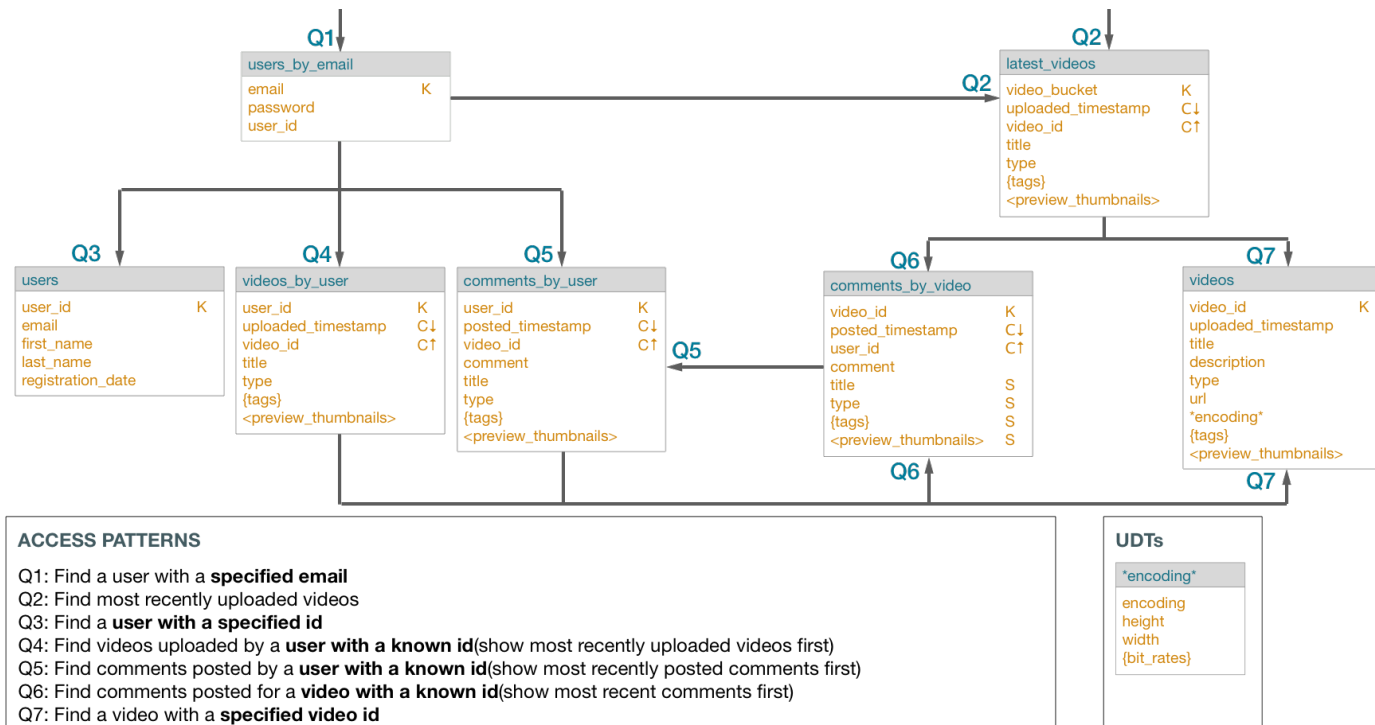


Physical UDT Diagram

- Represents user defined types and tuples



Example Chebotko Diagram

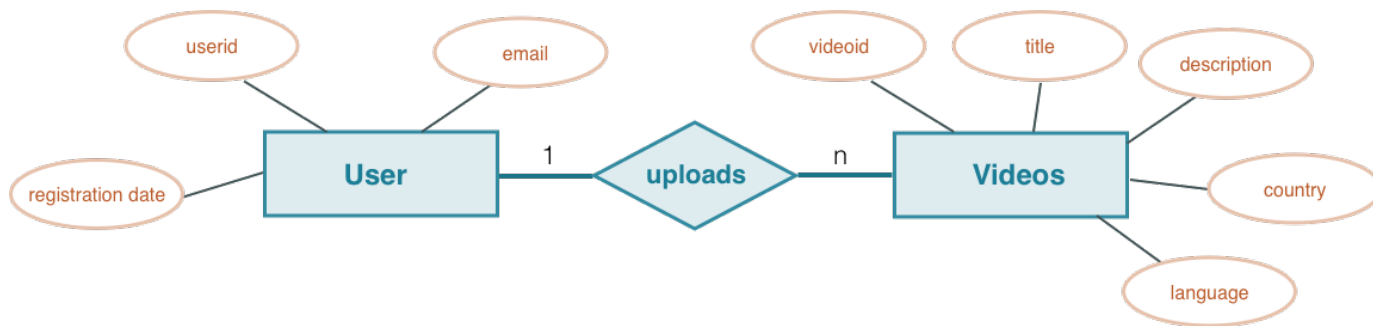


Data Modeling Principles

- Know your data
- Know your queries
- Nest data
- Duplicate data

Know Your Data

- Data captured by conceptual data model
- Define what is stored in database
- Preserve properties so that data is organized correctly



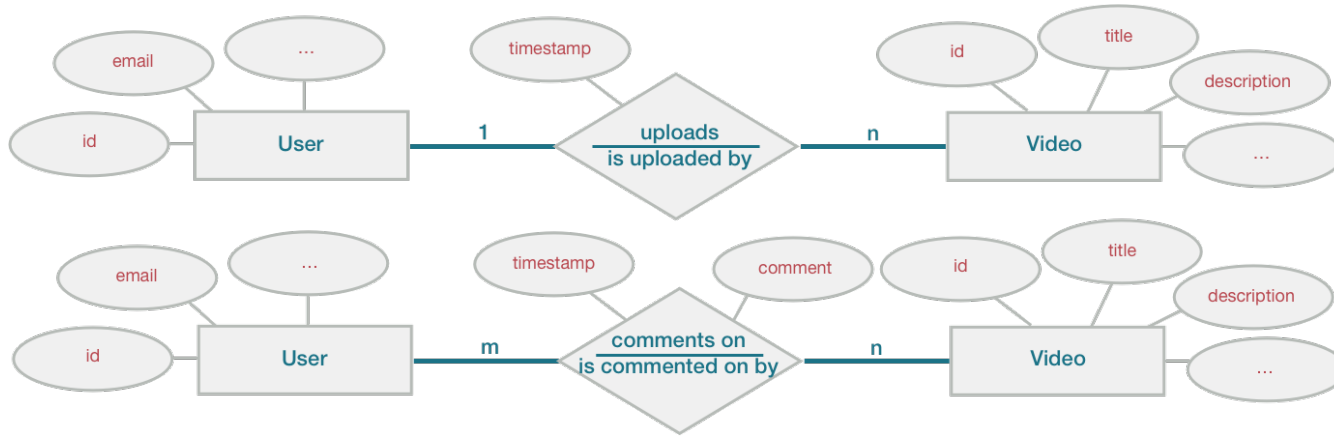
Know your Data

- Entity and relationship keys affect the table primary keys
- Primary key uniquely identifies a row / entity / relationship
- Composed of a key and possibly additional columns

videos	
video_id	K
uploaded_timestamp	
user_id	
title	
description	
type	
encoding	
{tags}	
<preview_thumbnails>	
{genres}	

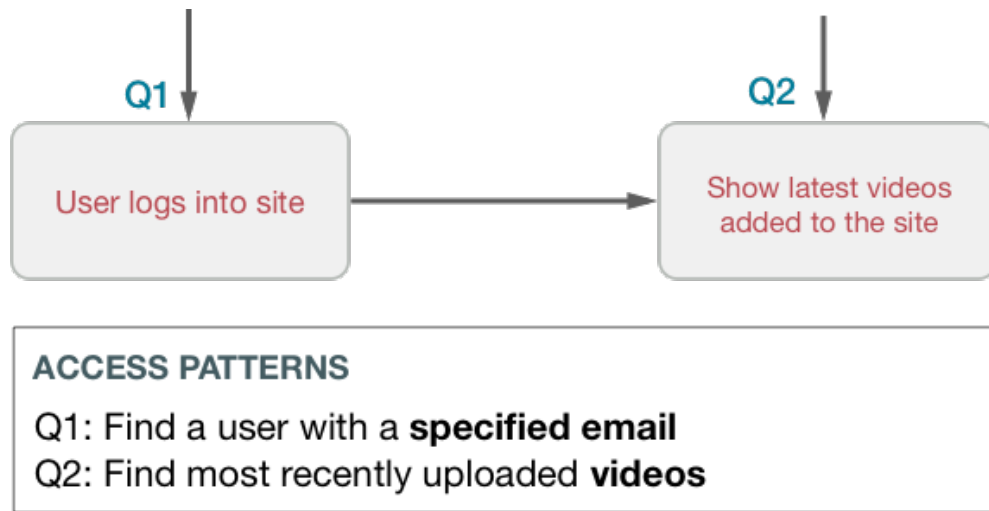
videos_by_user	
user_id	K
uploaded_timestamp	C↓
video_id	C↑
title	
type	
{tags}	
<preview_thumbnails>	

Cardinality Constraints Affect the Key for Relationships



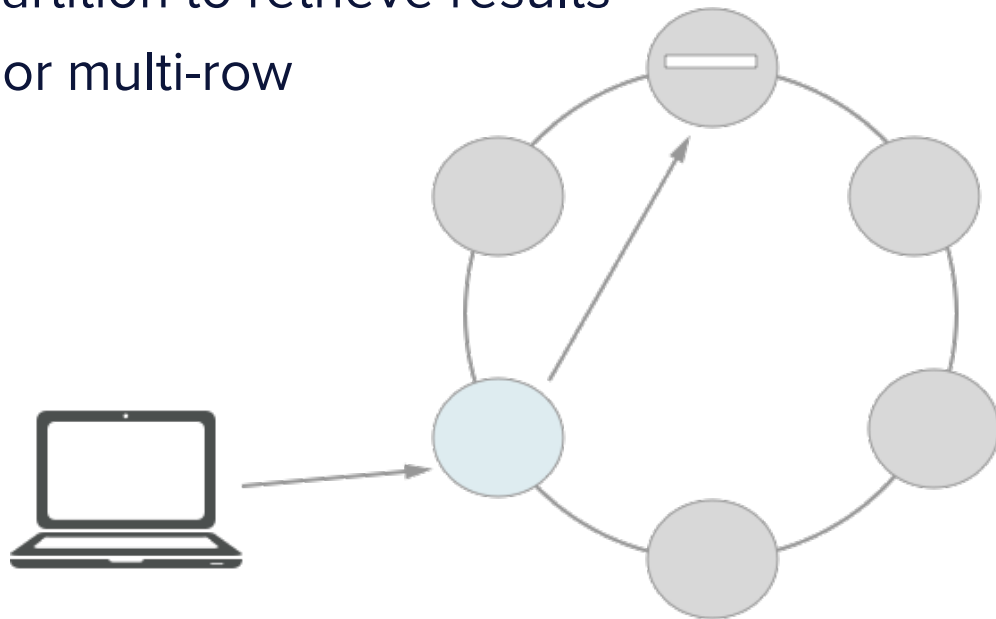
Know your Queries

- Queries captured by application workflow model
- Table schema design changes if queries change



Single Partition per Query = Ideal

- Most efficient access pattern
- Query accesses only one partition to retrieve results
- Partition can be single-row or multi-row



Partition+ Per Query = Acceptable

- Less efficient access pattern but not necessarily bad
- Query needs to access multiple partitions to retrieve results






Table Scan/Multi-Table Scan = Anti-Pattern


- Least efficient type of query but may be needed in some cases
- Query needs to access all partitions in a table(s) to retrieve results



Logical Model



Where Are We Now?



Nest Data

- Data nesting is the main data modeling technique
- Nesting organizes multiple entities into a single partition
- Supports partition per query data access
- Three data nesting mechanisms:
 - Clustering columns – multi-row partitions
 - Collection columns
 - User-defined type columns

Nest Data - Clustering Columns

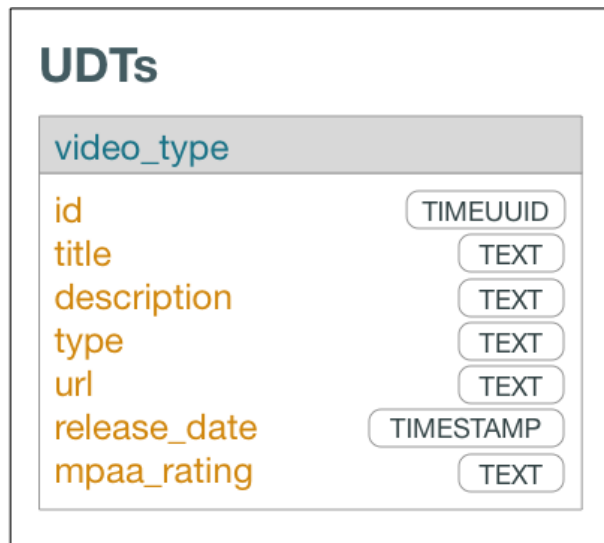

- Clustering column—primary data nesting mechanism
- Partition key identifies an entity that other entities will nest into
- Values in a clustering column identify the nested entities
- Multiple clustering columns implement multi-level nesting

videos	
video_id	K
uploaded_timestamp	
user_id	
title	
description	
type	
{tags}	
<preview_thumbnails>	
{genres}	

actors_by_video	
video_id	K
actor_name	C↑
character_name	C↑

Nest Data – UDT

- User-defined type—secondary data nesting mechanism
- Represents 1-1 relationship, but can use in conjunction with collections
- Easier than working with multiple collection columns




Duplicate Data

- Partition per query and data nesting may result in data duplication
- Query results are pre-computed and materialized
- Data can be duplicated across tables, partitions, and / or rows

videos_by_actor	videos_by_genre	videos_by_tag
actor K	genre K	tag K
release_date C↓	release_date C↓	release_date C↓
video_id C↑	video_id C↑	video_id C↑
title	title	title
type	type	type
{tags}	{tags}	{tags}
<preview_thumbnails>	<preview_thumbnails>	<preview_thumbnails>

Mapping Rules



Mapping Rules for Query Driven Methodology

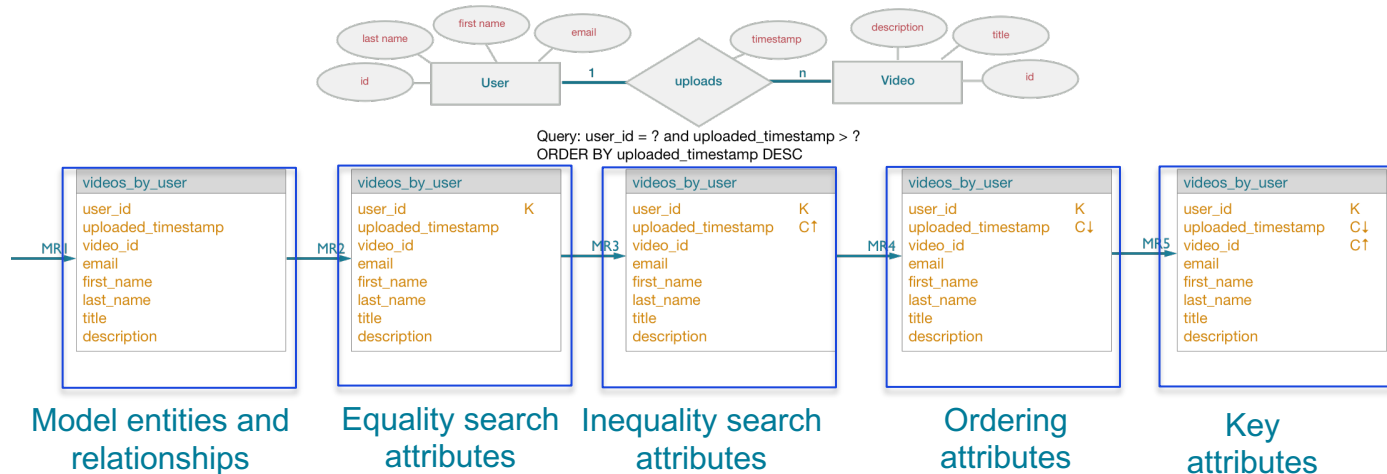
- Mapping rules ensure that a logical data model is correct
- Each query has a corresponding table
- Tables are designed to allow queries to execute properly
- Tables return data in the correct order

What are the Rules?

- Mapping Rule 1: Entities and relationships
- Mapping Rule 2: Equality search attributes
- Mapping Rule 3: Inequality search attributes
- Mapping Rule 4: Ordering attributes
- Mapping Rule 5: Key attributes

Mapping Rules in Action


- Create a table schema from the conceptual data model and for each query
- Apply the mapping rules in order



Exercise 03.03




Finalize your Logical Model




Exercise 03.03 – Finalize your Logical Model

- Add tables to the logical model to support additional queries
- Ensure that the mapping rules are applied appropriately

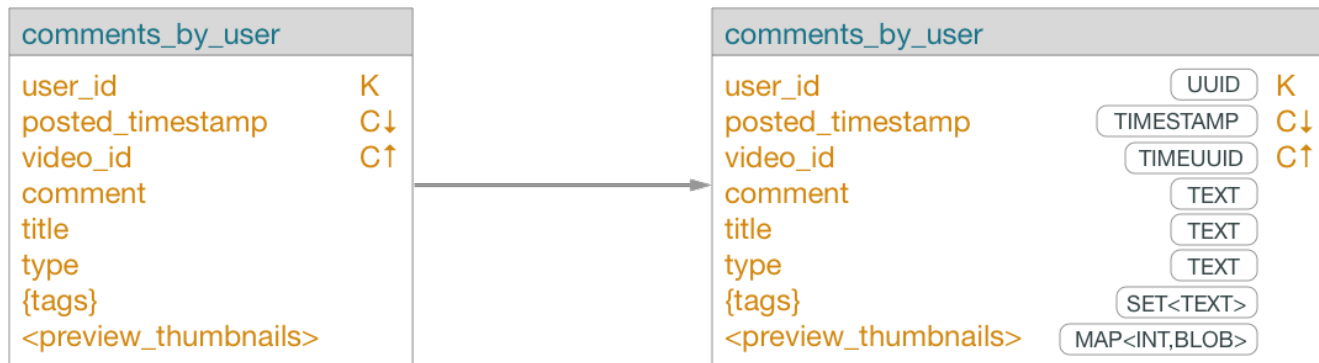
Physical Model



Where Are We Now?



Adding Data Types



Creating Tables

```
CREATE TABLE comments_by_user (  
    user_id UUID,  
    posted_timestamp TIMESTAMP,  
    video_id TIMEUUID,  
    comment TEXT,  
    title TEXT,  
    type TEXT,  
    tags SET<TEXT>,  
    preview_thumbnails MAP<INT, BLOB>,  
    PRIMARY KEY ((user_id), posted_timestamp, video_id)  
) WITH CLUSTERING ORDER BY (posted_timestamp DESC, video_id ASC);
```

Data Loading Methods

- COPY Command
- SSTable Loader
- Spark for Data Loading
- More about loading in Module 5!

Basic Data Loading with CQLCOPY


- COPY TO exports data from a table to a CSV file
- COPY FROM imports data to a table from a CSV file
- The process verifies the PRIMARY KEY and updates existing records
- If HEADER = false is specified the fields are imported in deterministic order
- When column names are specified, fields are imported in that order—missing and empty fields set to null
- Source cannot have more fields than the target table--can have fewer fields

```
COPY table1 (column1, column2, column3) FROM 'table1data.csv'  
WITH HEADER=true;
```

Exercise 03.04



Finalize your Physical Model



Exercise 03.04 – Finalize your Physical Model

- Add data types to the physical data model
- Run the **CQL CREATE TABLE** statements for each table in physical model
- Load data and run some queries to test the physical data model